

APPLICATION NOTE

lifelines-hc: a Python package for Higher Criticism testing in two-sample survival analysis under non-proportional hazards

Alon Kipnis,^{1,*} Ben Galili² and Zohar Yakhini^{1,2}¹School of Computer Science, Reichman University, Herzliya, 4610101, Israel and ²Department of Computer Science, Technion – Israel Institute of Technology, Haifa, 3200003, Israel

*E-mail: alon.kipnis@runi.ac.il

FOR PUBLISHER ONLY Received on Month ; revised on Month ; accepted on Month

Abstract

Motivation: The log-rank test is a standard tool for two-sample survival comparison and has good power against the widely studied proportional-hazards alternatives. However, Kipnis et al. [2026] showed that it can lose power against sparse hazard departures, where the hazard difference is concentrated in a small number of time intervals whose locations along the follow-up are unknown a priori. Weighted log-rank tests, such as Gehan-Wilcoxon, Tarone-Ware, Peto-Prentice, and Fleming-Harrington, each impose a pre-specified temporal emphasis. Combination procedures such as MaxCombo [Lin et al., 2020] and the Yang-Prentice short-term/long-term hazard-ratio model [Yang and Prentice, 2005, 2010] relax that choice, but still scan only a small dictionary of global temporal shapes, and are therefore also insensitive to sparse departures occurring outside those shapes. **Results:** We introduce **lifelines-hc**, a Python package extending the **lifelines** survival library with the HCHG test: Higher Criticism (HC) applied to per-interval HyperGeometric (HG) p-values from Kipnis et al. [2026]. HCHG is a test for right-censored two-sample survival data. It detects sparse hazard departures at unknown locations without committing to any temporal pattern. Across four biological case studies — an immuno-oncology trial (CheckMate 057 PFS, $n = 582$), a targeted kinase-inhibitor trial in metastatic prostate cancer (COMET-1 OS, $n = 1028$), an adjuvant bisphosphonate trial (AZURE DFS, $n = 3359$), and the Copenhagen Study Group for Liver Diseases (CSL) liver-cirrhosis trial with fully public individual patient data ($n = 446$) — HCHG achieves $p \leq 0.014$ in every dataset, while the log-rank test is non-significant throughout ($p \geq 0.26$). MaxCombo and the Yang-Prentice adaptive log-rank test detect the delayed-benefit crossing on CheckMate 057 ($p < 0.001$) but are non-significant on the remaining three datasets ($p \geq 0.24$), where the departure is sparse or multi-window rather than a smooth short-term/long-term hazard-ratio pattern. **Availability and Implementation:** **lifelines-hc** is freely available under the MIT license at <https://github.com/alonkipnis/lifelines-hc> and via `pip install lifelines-hc`; it requires Python ≥ 3.8 and **lifelines** ≥ 0.29 . Full analysis scripts and per-method results tables are provided as Supplementary Data.

Key words: survival analysis; higher criticism; sparse signal detection; hypergeometric test; hypothesis testing; Python

Introduction

The Kaplan-Meier estimator [Kaplan and Meier, 1958] and the log-rank test [Mantel, 1966] are the workhorses of two-sample survival analysis in biomedical research. The log-rank statistic has good power against a wide range of useful alternatives, particularly proportional-hazards (PH) alternatives [Schoenfeld, 1981]. Nevertheless, Kipnis et al. [2026] showed that the log-rank test is not optimal against *sparse* hazard departures, in which the hazard difference is concentrated in a small, unknown subset of time intervals.

Such localized or non-proportional departures arise naturally in several biomedical settings. Immune checkpoint inhibitors may require an immune-priming phase before producing durable tumour suppression, leading to delayed treatment effects and

crossing survival curves; for example, in the second-line non-squamous NSCLC trial CheckMate 057, nivolumab improved overall survival relative to docetaxel, while the progression-free survival curves crossed, consistent with delayed benefit [Borghaei et al., 2015]. Targeted therapies can also produce endpoint-specific benefits over a limited follow-up window that are diluted in global survival summaries: in the COMET-1 mCRPC trial, cabozantinib markedly improved bone scan response at week 12 and radiographic progression-free survival, but did not significantly improve overall survival [Smith et al., 2016]. Subgroup-dependent biology can similarly generate composite, time-varying hazard differences: in the AZURE trial of adjuvant zoledronic acid, no overall benefit was observed in the full population, but benefit was seen among women with

established menopause, consistent with an effect mediated by the postmenopausal bone microenvironment [Coleman et al., 2011, 2014]. Time-varying treatment effects can likewise produce multi-window hazard departures: in the CSL liver-cirrhosis trial, prednisone was harmful early (infection and fluid retention) but beneficial later through anti-inflammatory action, yielding survival curves that cross more than once [Martinussen and Scheike, 2006]. This canonical non-proportional-hazards dataset is widely used to illustrate time-varying-coefficient models and provides fully public individual patient data via the CRAN `timereg` package [Martinussen and Scheike, 2006] and the SurvSet repository [Drysdale, 2022].

The need for log-rank alternatives in survival analysis is well recognized, and several weighted log-rank tests have been developed for non-proportional-hazards settings [Gehan, 1965, Peto and Peto, 1972, Tarone and Ware, 1977, Harrington and Fleming, 1982, Fleming and Harrington, 1991]. Each applies a temporal weight function, emphasizing early, intermediate, or late events. Each of these tests achieves good power when the signal matches the intended temporal pattern. However, choosing an appropriate weight requires knowing in advance *where* the hazard departure will occur. When the departure is sparse but its location is unknown a priori, all weighted log-rank tests can miss it simultaneously, as we demonstrate below.

Contemporary practice therefore combines several Fleming–Harrington weights rather than committing to one. The MaxCombo test [Lin et al., 2020, Ristl et al., 2021] takes the maximum of $\text{FH}(\rho, \gamma)$ statistics over the dictionary $(\rho, \gamma) \in \{(0, 0), (0, 1), (1, 0), (1, 1)\}$ — log-rank, late, early, and middle emphasis — and obtains a multiplicity-adjusted p -value from the estimated multivariate normal joint distribution of the component Z -statistics. The Yang–Prentice short-term/long-term hazard-ratio model [Yang and Prentice, 2005] interpolates between an immediate hazard ratio θ_1 and an asymptotic ratio θ_2 , covering delayed, diminishing, and crossing hazards with two parameters; the associated adaptive weighted log-rank test [Yang and Prentice, 2010] is a leading model-based competitor. Both procedures are powerful against the global early-versus-late patterns they were designed for, particularly delayed benefit in immuno-oncology. They remain, however, averaging methods over a small dictionary of global temporal shapes, and can miss sparse or multi-window departures whose locations are unknown a priori.

Kipnis et al. [2026] proposed the HCHG test — the Higher Criticism (HC) statistic of Donoho and Jin [2004] applied to per-interval HyperGeometric (HG) p -values — specifically for this setting. HCHG divides the follow-up into equal-width intervals, derives a hypergeometric p -value for the event-count comparison in each interval under the null of equal group-specific hazards, and applies the HC statistic to test whether the body of these p -values is collectively implausible under the global null. Because HCHG does not pre-specify any temporal weight, it retains power for sparse departures wherever they occur along the follow-up axis.

We present `lifelines-hc`, a Python package that implements HCHG and integrates with the popular `lifelines` library [Davidson-Pilon, 2019].

The lifelines-hc package

The HCHG statistic

HCHG [Kipnis et al., 2026] combines two components: hypergeometric interval p -values and the HC statistic of Donoho and Jin [2004, 2015].

Hypergeometric interval p -values. Given two groups with right-censored observations (T_i, E_i, G_i) , partition the follow-up range into K equal-width intervals. For interval k , let n_k be the total events, r_k^A and r_k^B the subjects at risk in groups A and B, and x_k the events in group B. Under the null of equal group-specific hazards,

$$x_k \mid n_k, r_k^A, r_k^B \sim \text{Hypergeometric}(r_k^A + r_k^B, r_k^B, n_k). \quad (1)$$

This yields a one-sided p -value p_k for each interval using an exact test.

HC aggregation. Ordering the K p -values as $p_{(1)} \leq \dots \leq p_{(K)}$, the HC statistic [Donoho and Jin, 2004, 2015] is applied to this collection:

$$\text{HCHG}_K = \max_{1 \leq j \leq \lfloor \gamma K \rfloor} \frac{j/K - p_{(j)}}{\sqrt{p_{(j)}(1 - p_{(j)})/K}}, \quad (2)$$

where $\gamma \in (0, 1)$ (default 0.2) restricts scanning to the most extreme fraction of interval p -values [Donoho and Jin, 2004, 2015]. A global p -value is obtained by permuting group labels ($n_{\text{perms}} = 2000$ in the analyses reported here). For the two-sided alternative, the minimum of the two one-sided interval p -values replaces p_k in (2).

Software design and API

`lifelines-hc` exposes a function-based API consistent with `lifelines.statistics`, accepting NumPy arrays of event times and indicators:

```
from lifelines_hc import (higher_criticism_test,
                          suspected_deviations)

result = higher_criticism_test(
    T_A, T_B,
    event_observed_A=E_A, event_observed_B=E_B,
    n_intervals_to_pool=100, n_permutations=2000,
    alternative='both', gamma=0.2, seed=42)

print(result.p_value)           # permutation-
                                # calibrated p-value
print(result.test_statistic)    # HCHG statistic
                                # value
```

The companion `suspected.deviations()` function returns a `pandas.DataFrame` of per-interval p -values with a `suspected` boolean column marking HCHG-flagged intervals, enabling overlay of diagnostic shading on Kaplan-Meier plots. The package additionally ships with plotting utilities (`plot_km_with_hc`, `plot_pvalue_profile`) and a convenience wrapper `run_all_tests()` that benchmarks HCHG, log-rank, the four weighted alternatives, MaxCombo, and Yang–Prentice in a single call, returning a `pandas.DataFrame` for easy comparison.

Results

We demonstrate `lifelines-hc` on four biological datasets with distinct non-proportional-hazards (PH) structures (Figure 1),

benchmarking HCHG against the log-rank test, four weighted alternatives, the MaxCombo combination test, and the Yang–Prentice adaptive log-rank test. Individual patient data for the first three trials were reconstructed from published Kaplan–Meier curves; the CSL trial provides publicly available individual patient data. Full per-method p -value tables are provided in Supplementary Data.

Immuno-oncology: crossing survival curves

CheckMate 057 [Borghaei et al., 2015] randomised 582 patients with advanced non-squamous NSCLC to nivolumab ($n = 292$) or docetaxel ($n = 290$). Progression-free survival showed crossing curves: docetaxel provided a short-term advantage during immune priming, subsequently overtaken by the durable nivolumab response (Figure 1A). Individual patient data were reconstructed from the published Kaplan–Meier curve via the algorithm of Guyot et al. [2012] as implemented in the `kmdata` R package.

The log-rank test is non-significant ($p = 0.351$) as the early excess and late benefit partially cancel. Among the weighted alternatives, Gehan–Wilcoxon ($p = 0.041$) and Peto–Prentice ($p = 0.049$) reach borderline significance by detecting the early crossing, but Tarone–Ware ($p = 0.358$) and Fleming–Harrington (1,1) ($p = 0.256$) do not. This is the delayed-benefit pattern MaxCombo and Yang–Prentice are designed for: MaxCombo is highly significant ($p = 2.5 \times 10^{-5}$), driven by the late-weighted FH(0,1) member ($Z = 4.27$), and the Yang–Prentice adaptive log-rank test is likewise significant ($p = 0.00055$), with estimated short-term hazard ratio $\hat{\theta}_1 = 2.33$ (nivolumab excess during priming) and long-term ratio $\hat{\theta}_2 = 0.44$ (durable nivolumab benefit). HCHG achieves $p = 0.001$ and characterises the full crossing structure by flagging both the early divergence and late separation windows (Figure 1A,E).

Targeted therapy in prostate cancer: a transient early benefit

The COMET-1 trial [Smith et al., 2016] randomised 1028 patients with chemotherapy-pretreated mCRPC to cabozantinib ($n = 682$) or prednisone ($n = 346$) in a 2:1 ratio. Overall survival showed no significant log-rank difference ($p = 0.262$). Cabozantinib, by inhibiting MET and VEGFR2 kinases, significantly improved bone scan response at week 12 (a pre-specified secondary endpoint), confirming genuine biological activity against bone-metastatic disease. This short-term activity did not, however, translate into a sustained overall survival (OS) benefit. The result is a temporally concentrated early OS advantage (months 3–7) that is subsequently lost (Figure 1B,F).

All four weighted alternatives are non-significant ($p \geq 0.19$): the early-weighted Gehan–Wilcoxon ($p = 0.193$) and Peto–Prentice ($p = 0.206$) come closest but are diluted by the null period after the bone-response window. MaxCombo ($p = 0.338$) and Yang–Prentice ($p = 0.262$) likewise miss the signal: a transient few-month advantage is not a global early-versus-late hazard-ratio shape. HCHG ($p = 0.014$) aggregates evidence from the specific early intervals where the transient effect is concentrated.

Adjuvant bisphosphonate therapy: a menopause-dependent effect

The AZURE trial [Coleman et al., 2011, 2014] randomised 3359 early-stage breast cancer patients to standard adjuvant therapy

with ($n = 1681$) or without ($n = 1678$) zoledronic acid. Disease-free survival showed no significant log-rank difference ($p = 0.305$) because the drug’s bone-microenvironment benefit depends on the low-oestrogen state of postmenopause, which was absent in the many premenopausal patients enrolled at baseline. As these patients progressively transitioned to postmenopause during the decade-long follow-up, the hazard difference accumulated in a specific mid-to-late window (months 20–60) rather than uniformly (Figure 1C,G).

All four weighted alternatives are non-significant ($p \geq 0.28$): no single temporal weight captures a pattern distributed across this window while remaining unaffected by the null early period. MaxCombo ($p = 0.401$) and Yang–Prentice ($p = 0.308$) are similarly non-significant. HCHG ($p = 0.005$) detects this scattered but collectively significant signal. The result is validated by published subgroup analyses showing significant disease-free survival (DFS) benefit in postmenopausal patients [Coleman et al., 2011, 2014], confirming that HCHG is detecting a genuine biological effect.

Corticosteroid therapy in liver cirrhosis: a time-varying effect

The CSL trial [Martinussen and Scheike, 2006] randomised 446 patients with histologically verified liver cirrhosis to prednisone ($n = 226$) or placebo ($n = 220$) in Copenhagen (1962–69), with 270 deaths over follow-up of up to ~ 13 years. Unlike the preceding case studies, the raw individual patient data are genuinely public—distributed in the CRAN `timereg` package (dataset `cs1`) and the open SurvSet repository [Drysdale, 2022]—rather than reconstructed from a published Kaplan–Meier curve.

Corticosteroids exert a time-varying effect in cirrhosis: early harm through immunosuppression, infection, and fluid retention is followed by later anti-inflammatory benefit in actively inflamed subgroups. The survival curves cross more than once, with the treatment effect concentrated in a few non-contiguous windows (Figure 1D,H).

The log-rank test is non-significant ($p = 0.384$), and all four weighted alternatives are as well ($p \geq 0.13$): Gehan–Wilcoxon ($p = 0.51$), Tarone–Ware ($p = 0.38$), Peto–Prentice ($p = 0.45$), and Fleming–Harrington (1,1) ($p = 0.14$). MaxCombo ($p = 0.236$) and Yang–Prentice ($p = 0.393$) also fail: no single temporal weight or two-parameter short-term/long-term hazard ratio captures a multi-window, sign-changing departure. HCHG achieves $p = 0.002$, demonstrating detection of sparse non-proportional-hazards departures on genuinely public raw data.

To summarise, across the four case studies, HCHG achieves $p \leq 0.014$ in every domain, while the log-rank test is non-significant throughout ($p \geq 0.26$). MaxCombo and Yang–Prentice are significant only on CheckMate 057, the delayed-benefit crossing they were designed to detect; on COMET-1, AZURE, and CSL they return $p \geq 0.24$. The two borderline significances for Gehan–Wilcoxon ($p = 0.041$) and Peto–Prentice ($p = 0.049$) in CheckMate 057 are the only other instances where any weighted alternative reaches $p < 0.05$. HCHG is therefore complementary to MaxCombo and Yang–Prentice: it matches them on global early-versus-late non-proportionality and retains power for sparse or multi-window departures that a small Fleming–Harrington dictionary cannot represent.

Conclusion

lifelines-hc provides an accessible Python implementation of HCHG for two-sample right-censored survival data. HCHG is specifically designed for sparse hazard departures at unknown temporal locations: unlike weighted log-rank alternatives, MaxCombo, and the Yang–Prentice adaptive test, it requires no a priori specification of where the hazard difference will occur and retains power for any concentrated departure along the follow-up axis. The package integrates transparently with **lifelines** workflows and returns diagnostic interval-level information identifying which time windows drive the signal.

Conflict of Interest

None declared.

Funding

The work of Alon Kipnis was supported in part by the US-Israel Binational Science Foundation (BSF) grant 2022124.

Data availability

Individual patient data for CheckMate 057, COMET-1, and AZURE were reconstructed from published Kaplan-Meier curves using the algorithm of Guyot et al. [2012] as implemented in the **kmdata** R package (<https://github.com/raredd/kmdata>). CSL liver-cirrhosis individual patient data are fully public via the CRAN **timereg** package (dataset **cs1**) and the **SurvSet** repository [Drysdale, 2022] (<https://github.com/ErikinBC/SurvSet>). All analysis code is available at <https://github.com/alonkipnis/lifelines-hc/tree/main/AppNote>.

References

- Hossein Borghaei, Luis Paz-Ares, Leora Horn, David R. Spigel, Martin Steins, Neal E. Ready, Laura Q. Chow, Everett E. Vokes, Enriqueta Felip, Esther Holgado, et al. Nivolumab versus docetaxel in advanced nonsquamous non-small-cell lung cancer. *New England Journal of Medicine*, 373(17): 1627–1639, 2015. doi: 10.1056/NEJMoa1507643.
- Robert E. Coleman, Howard Marshall, David Cameron, David Dodwell, Robert Burkinshaw, Mary Keane, Michael Gil, Stephen J. Houston, Robert J. Grieve, Peter J. Barrett-Lee, et al. Breast-cancer adjuvant therapy with zoledronic acid. *New England Journal of Medicine*, 365(15):1396–1405, 2011. doi: 10.1056/NEJMoa1105195.
- Robert E. Coleman, Marie Collinson, Eleanor Banks, Graham Clack, Dorothy Ritchie, Victoria Liversedge, David Cameron, David Dodwell, Michael Jones, Peter Barrett-Lee, et al. Benefits and risks of adjuvant treatment with zoledronic acid in stage II/III breast cancer. 10 year follow-up of the AZURE randomized clinical trial. *Journal of Bone Oncology*, 3(2): 23–30, 2014. doi: 10.1016/j.jbo.2013.12.001.
- Cameron Davidson-Pilon. **lifelines**: survival analysis in Python. *Journal of Open Source Software*, 4(40):1317, 2019. doi: 10.21105/joss.01317.
- David Donoho and Jiashun Jin. Higher criticism for detecting sparse heterogeneous mixtures. *The Annals of Statistics*, 32(3):962–994, 2004. doi: 10.1214/009053604000000265.
- David L. Donoho and Jiashun Jin. Higher criticism for large-scale inference, especially for rare and weak effects. *Statistical science*, 30(1):1–25, 2015.
- Erik Drysdale. **SurvSet**: An open-source time-to-event dataset repository. *arXiv preprint arXiv:2203.03094*, 2022. doi: 10.48550/arXiv.2203.03094.
- Thomas R. Fleming and David P. Harrington. *Counting Processes and Survival Analysis*. Wiley, New York, 1991. ISBN 0-471-52218-X.
- Edmund A. Gehan. A generalized Wilcoxon test for comparing arbitrarily singly-censored samples. *Biometrika*, 52(1–2): 203–223, 1965. doi: 10.1093/biomet/52.1-2.203.
- Patricia Guyot, A. E. Ades, Mario J. N. M. Ouwers, and Nicky J. Welton. Enhanced secondary analysis of survival data: reconstructing the data from published Kaplan-Meier survival curves. *BMC Medical Research Methodology*, 12:9, 2012. doi: 10.1186/1471-2288-12-9.
- David P. Harrington and Thomas R. Fleming. A class of rank test procedures for censored survival data. *Biometrika*, 69(3): 553–566, 1982. doi: 10.1093/biomet/69.3.553.
- Edward L. Kaplan and Paul Meier. Nonparametric estimation from incomplete observations. *Journal of the American Statistical Association*, 53(282):457–481, 1958. doi: 10.1080/01621459.1958.10501452.
- A Kipnis, B Galili, and Z Yakhini. Higher criticism for rare and weak non-proportional hazard deviations in survival analysis. *Biometrika*, 113(1):asaf075, 02 2026. ISSN 1464-3510. doi: 10.1093/biomet/asaf075. URL <https://doi.org/10.1093/biomet/asaf075>.
- Ray S. Lin, Ji Lin, Satrajit Roychoudhury, Keaven M. Anderson, Tianle Hu, Bo Huang, Larry F. Leon, Jason J. Z. Liao, Rong Liu, Xiaodong Luo, Pralay Mukhopadhyay, Rui Qin, Kay Tatsuoka, Xuejing Wang, Yang Wang, Jian Zhu, Tai-Tsang Chen, Renee Iacona, and Cross-Pharma Non-proportional Hazards Working Group. Alternative analysis methods for time to event endpoints under nonproportional hazards: a comparative analysis. *Statistics in Biopharmaceutical Research*, 12(2):187–198, 2020. doi: 10.1080/19466315.2019.1697738.
- Nathan Mantel. Evaluation of survival data and two new rank order statistics arising in its consideration. *Cancer Chemotherapy Reports*, 50(3):163–170, 1966.
- Torben Martinussen and Thomas H. Scheike. *Dynamic Regression Models for Survival Data*. Springer, New York, 2006. doi: 10.1007/0-387-33960-4. CSL liver cirrhosis trial (prednisone vs placebo, Copenhagen Study group for Liver diseases, 1962–1969); dataset **cs1** in the CRAN **timereg** package.
- Richard Peto and Julian Peto. Asymptotically efficient rank invariant test procedures. *Journal of the Royal Statistical Society, Series A*, 135(2):185–207, 1972. doi: 10.2307/2344317.
- Robin Ristl, Nicolas Ballarini, Heiko Götte, Armin Schüler, Martin Posch, and Franz König. Delayed treatment effects, treatment switching and heterogeneous patient populations: how to design and analyze RCTs in oncology. *Pharmaceutical Statistics*, 20(1):129–145, 2021. doi: 10.1002/pst.2062.
- David Schoenfeld. The asymptotic properties of nonparametric tests for comparing survival distributions. *Biometrika*, 68(1): 316–319, 1981. doi: 10.1093/biomet/68.1.316.
- Matthew R. Smith, Christopher J. Sweeney, Paul G. Corn, Dana E. Rathkopf, David C. Smith, Maha Hussain, Daniel J. George, Celestia S. Higano, Andrea L. Harzstark, A. Oliver Sartor, et al. Cabozantinib in chemotherapy-pretreated




Fig. 1. Higher Criticism detects non-proportional hazard departures across four clinical datasets with distinct temporal patterns. Top row (A–D): Kaplan-Meier curves with HCHG-flagged intervals (shaded). Bottom row (E–H): per-interval $\text{signed } -\log_{10}(\text{hypergeometric } p\text{-value})$. Upward bars mark excess events in the treatment arm, downward bars mark excess in the control arm. Dashed lines are the per-direction HCHG thresholds; bars reaching beyond either line are highlighted in red and coincide exactly with the intervals shaded above. (A,E) CheckMate 057 PFS (nivolumab vs. docetaxel, $n = 582$); log-rank $p = 0.351$, HCHG $p = 0.001$. Curves cross at ~ 3 months before nivolumab gains a durable late advantage. (B,F) COMET-1 OS (cabozantinib vs. prednisone, $n = 1028$); log-rank $p = 0.262$, HCHG $p = 0.014$. Early OS advantage (months 3–7) during bone-lesion response, fading as resistance develops. (C,G) AZURE DFS (zoledronic acid vs. control, $n = 3359$); log-rank $p = 0.305$, HCHG $p = 0.005$. Menopause-dependent benefit accumulates in months 20–60. (D,H) CSL OS (prednisone vs. placebo, liver cirrhosis, $n = 446$); log-rank $p = 0.384$, HCHG $p = 0.002$. Time-varying corticosteroid effect with multiple curve crossings; *fully public* raw individual patient data (CRAN `timereg` / `SurvSet`). In every dataset the log-rank test is non-significant while HCHG is significant. MaxCombo and Yang–Prentice are significant only on the delayed-benefit crossing in (A,E); see Supplementary Data. For the first three, individual patient data were reconstructed via Guyot et al. [2012].

metastatic castration-resistant prostate cancer: results of a phase III randomized controlled trial (COMET-1). *Journal of Clinical Oncology*, 34(25):3005–3013, 2016. doi: 10.1200/JCO.2015.65.5597.

Robert E. Tarone and James Ware. On distribution-free tests for equality of survival distributions. *Biometrika*, 64(1):156–160, 1977. doi: 10.1093/biomet/64.1.156.

Song Yang and Ross Prentice. Semiparametric analysis of short-term and long-term hazard ratios with two-sample survival data. *Biometrika*, 92(1):1–17, 2005. doi: 10.1093/biomet/92.1.1.

Song Yang and Ross Prentice. Improved logrank-type tests for survival data using adaptive weights. *Biometrics*, 66(1): 30–38, 2010.